

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

Importing the Data

In [2]:

```
Df=pd.read_csv('health care diabetes.csv')
Df.head()
```

Out[2]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

In [3]:

```
Df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [4]:

```
Df['Outcome'].value_counts()
```

Out[4]:

```
0    500
1    268
Name: Outcome, dtype: int64
```

Histogram to see the distribution of the data

In [5]:

```
plt.figure(figsize=(30,30))
plt.subplot(3,2,1)
Df['BloodPressure'].plot(kind='hist',bins=100)
```

```
plt.title('BP', fontsize=25)

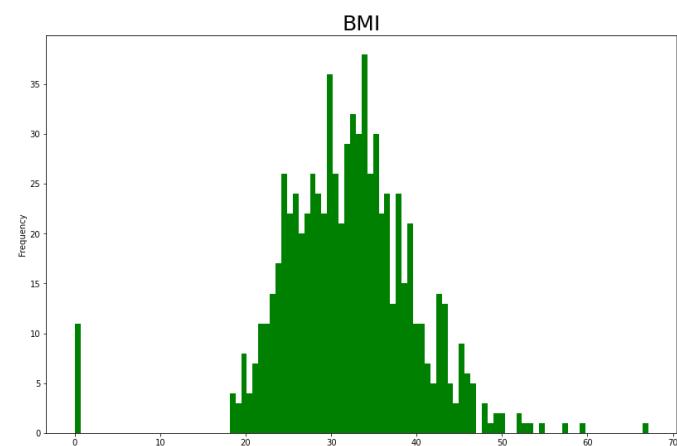

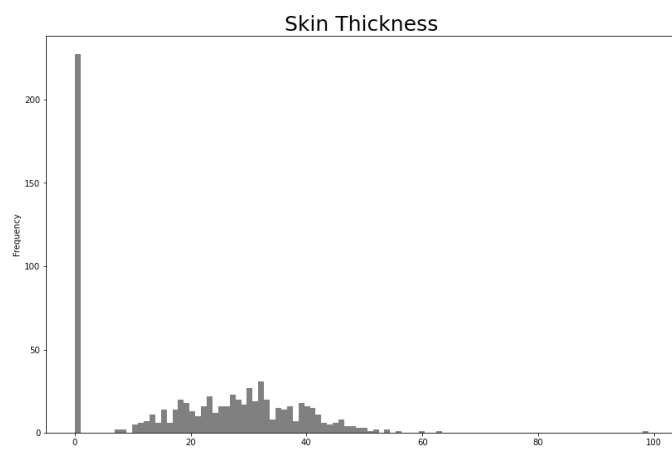

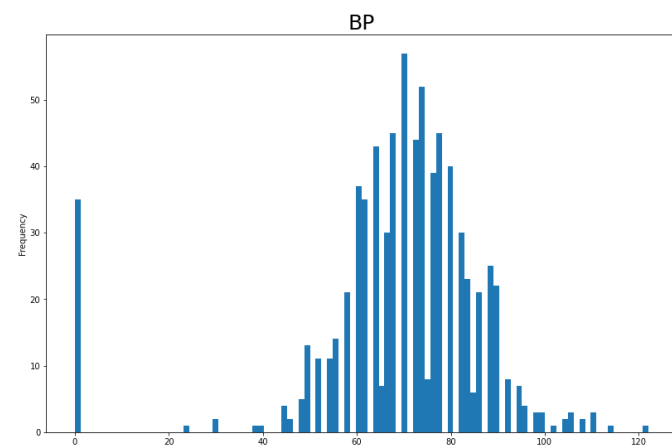
plt.subplot(3,2,2)
Df['Glucose'].plot(kind='hist', bins=100, color='Purple')
plt.title('Glucose', fontsize=25)

plt.subplot(3,2,3)
Df['SkinThickness'].plot(kind='hist', bins=100, color='Grey')
plt.title('Skin Thickness', fontsize=25)

plt.subplot(3,2,4)
Df['Insulin'].plot(kind='hist', bins=100, color='orange')
plt.title('Insulin', fontsize=25)

plt.subplot(3,2,5)
Df['BMI'].plot(kind='hist', bins=100, color='green')
plt.title('BMI', fontsize=25)

plt.show()
```



Finding the Average of each column

In [6]:

```
Df[['BloodPressure']].max()
```

```

BpAvg=Df[ 'BloodPressure' ].mean()
GlucoseAvg=Df[ 'Glucose' ].mean()
SkinThicknessAvg=Df[ 'SkinThickness' ].mean()
InsulinAvg=Df[ 'Insulin' ].mean()
BMIAvg=Df[ 'BMI' ].mean()

print(BpAvg)
print(GlucoseAvg)
print(SkinThicknessAvg)
print(InsulinAvg)
print(BMIAvg)

```

```

69.10546875
120.89453125
20.536458333333332
79.79947916666667
31.992578124999977

```

Converting Int and Float to Object for replacing the '0' values with the mean.

In [7]:

```

Df[ 'BloodPressure' ]=Df[ 'BloodPressure' ].map(str)
Df[ 'Glucose' ]=Df[ 'Glucose' ].map(str)
Df[ 'SkinThickness' ]=Df[ 'SkinThickness' ].map(str)
Df[ 'Insulin' ]=Df[ 'Insulin' ].map(str)
Df[ 'BMI' ]=Df[ 'BMI' ].map(str)

```

In [8]:

```
Df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Pregnancies           768 non-null    int64
 1   Glucose               768 non-null    object
 2   BloodPressure         768 non-null    object
 3   SkinThickness         768 non-null    object
 4   Insulin               768 non-null    object
 5   BMI                   768 non-null    object
 6   DiabetesPedigreeFunction 768 non-null    float64
 7   Age                  768 non-null    int64
 8   Outcome              768 non-null    int64
dtypes: float64(1), int64(3), object(5)
memory usage: 54.1+ KB

```

Replacing '0' with the Average

In [9]:

```

Df[ 'BloodPressure' ]=Df[ 'BloodPressure' ].replace('0',BpAvg)
Df[ 'Glucose' ]=Df[ 'Glucose' ].replace('0',GlucoseAvg)
Df[ 'SkinThickness' ]=Df[ 'SkinThickness' ].replace('0',SkinThicknessAvg)
Df[ 'Insulin' ]=Df[ 'Insulin' ].replace('0',InsulinAvg)
Df[ 'BMI' ]=Df[ 'BMI' ].replace('0',BMIAvg)

```

Converting again Object type to Int & Float for Model training

In [10]:

```

Df[ 'BloodPressure' ]=Df[ 'BloodPressure' ].map(int)
Df[ 'Glucose' ]=Df[ 'Glucose' ].map(int)
Df[ 'SkinThickness' ]=Df[ 'SkinThickness' ].map(int)
Df[ 'Insulin' ]=Df[ 'Insulin' ].map(int)
Df[ 'BMI' ]=Df[ 'BMI' ].map(float)

```

In [11]:

```
Df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null    int64
1   Glucose               768 non-null    int64
2   BloodPressure         768 non-null    int64
3   SkinThickness         768 non-null    int64
4   Insulin               768 non-null    int64
5   BMI                  768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                  768 non-null    int64
8   Outcome              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

Correlation among the variables

In [12]:

```
correlation=Df.corr()
correlation
```

Out[12]:


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
Pregnancies	1.000000	0.128022	0.208987	0.009393	0.018780	0.017683	-0.03352
Glucose	0.128022	1.000000	0.219765	0.158060	0.396137	0.235123	0.13715
BloodPressure	0.208987	0.219765	1.000000	0.130403	0.010492	0.242947	0.00047
SkinThickness	0.009393	0.158060	0.130403	1.000000	0.245410	0.499700	0.15719
Insulin	-0.018780	0.396137	0.010492	0.245410	1.000000	0.189561	0.15824
BMI	0.017683	0.235123	0.242947	0.499700	0.189561	1.000000	0.14064
DiabetesPedigreeFunction	-0.033523	0.137158	0.000471	0.157196	0.158243	0.140647	1.00000
Age	0.544341	0.266673	0.326791	0.020582	0.037676	0.036242	0.03356
Outcome	0.221898	0.492884	0.162879	0.171857	0.178696	0.292695	0.17384

In [13]:

```
plt.figure(figsize=(20,10))
plt.scatter(Df['SkinThickness'],Df['BMI'])
plt.title('Scatter Plot')
plt.xlabel('SkinThickness')
plt.ylabel('BMI')
```

Out[13]:

Text(0, 0.5, 'BMI')



Model Building

In [15]:

```
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
from sklearn import metrics
```

In [16]:

```
X=Df.drop('Outcome',axis=1)
y=Df['Outcome']
```

In [17]:

```
Xtrain,Xtest,ytrain,ytest=train_test_split(X,y,test_size=0.3,random_state=15)
```

In [18]:

```
svc=SVC()
```

In [19]:

```
svc.fit(Xtrain,ytrain)
```

Out[19]:

```
SVC()
```

In [20]:

```
ypredicted=svc.predict(Xtest)
```

In [21]:

```
metrics.confusion_matrix(ytest,ypredicted)
```

Out[21]:

```
array([[142,  18],
       [ 35,  36]], dtype=int64)
```

In [22]:

```
pd.crosstab(ytest, ypredicted, rownames=['True'], colnames=['Predicted'], margins=True)
```

Out[22]:

Predicted	0	1	All
True			
0	142	18	160
1	35	36	71
All	177	54	231

In [23]:

```
metrics.accuracy_score(ytest,ypredicted)
```

Out[23]:

0.7705627705627706

In [24]:

```
print(classification_report(ytest, ypredicted))
```

	precision	recall	f1-score	support
0	0.80	0.89	0.84	160
1	0.67	0.51	0.58	71
accuracy			0.77	231
macro avg	0.73	0.70	0.71	231
weighted avg	0.76	0.77	0.76	231

We can conclude that our model test accuracy is not bad. Now we will check if the model is generalised or not.

In [25]:

```
ytrainpred=svc.predict(Xtrain)
```

In [26]:

```
metrics.accuracy_score(ytrain,ytrainpred)
```

Out[26]:

0.7616387337057728


From the above test we saw that the model is working fine and has got Low Bias and Low Variance.

In [27]:

```
fpr, tpr, _ = metrics.roc_curve(ytest, ypredicted)
```

In [28]:

```
plt.plot(fpr,tpr)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



ROC curve shows the best threshold value for the classification

Lets try to build the Logistic Regression model and check whether it will give more accuracy or not

In [29]:

```
from sklearn.linear_model import LogisticRegression
```

In [30]:

```
lr=LogisticRegression()
```

In [31]:

```
lr.fit(Xtrain,ytrain)
```

C:\Users\NAVAL\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Out[31]:

```
LogisticRegression()
```

In [32]:

```
ypred=lr.predict(Xtest)
```

In [33]:

```
metrics.accuracy_score(ytest,ypred)
```

Out[33]:

```
0.7575757575757576
```

In [34]:

```
metrics.confusion_matrix(ytest,ypred)
```

Out[34]:

```
array([[136,  24],
       [ 32,  39]], dtype=int64)
```

In [35]:

```
pd.crosstab(ytest, ypred, rownames=['True'], colnames=['Predicted'], margins=True)
```


Out[35]:


Predicted	0	1	All
True			
0	136	24	160
1	32	39	71
All	168	63	231

In [36]:

```
fpr, tpr, _ = metrics.roc_curve(ytest, ypred)
plt.plot(fpr,tpr)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

10





We will go for the SVM model as it gave us more accuracy

KMeans Clustering

In [37]:

```
from sklearn.cluster import KMeans
```

In [38]:


```
wcss=[]
for i in range(1,11):
    kmeans=KMeans(n_clusters=i,init='k-means++',random_state=0)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)

plt.plot(range(1,11),wcss)
plt.title('Elbow Method')
plt.xlabel('No. of clusters')
plt.ylabel('WCSS')
plt.xticks(range(1, 11))
plt.show()
```

C:\Users\NAVAL\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:881: UserWarning: K Means is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=3

·

warnings.warn(



With the help of elbow method we concluded that we should use 2 clusters

In [39]:

```
clusters=KMeans(n_clusters=2,random_state=0)
```

In [40]:

```
clusters.fit(Xtrain,ytrain)
```

Out[40]:

```
KMeans(n_clusters=2, random_state=0)
```

In [41]:

```
yknn=clusters.predict(Xtest)
```

In [42]:

```
metrics.accuracy_score(ytest,yknn)
```

Out[42]:

```
0.7186147186147186
```

In [43]:

```
metrics.confusion_matrix(ytest,yknn)
```

Out[43]:

```
array([[152,   8],
       [ 57,  14]], dtype=int64)
```

In [44]:

```
pd.crosstab(ytest, yknn, rownames=['True'], colnames=['Predicted'], margins=True)
```

Out[44]:

Predicted	0	1	All
True			
0	152	8	160
1	57	14	71
All	209	22	231

We are getting higher TN but lower TP from KNN clustering as compared to the SVM model

In [45]:

```
print(classification_report(ytest, yknn))
```


	precision	recall	f1-score	support
0	0.73	0.95	0.82	160
1	0.64	0.20	0.30	71
accuracy			0.72	231
macro avg	0.68	0.57	0.56	231
weighted avg	0.70	0.72	0.66	231

In [46]:


```
Df.to_csv(r'C:\Users\NAVAL\Downloads\Healthcare_diabetes.csv')
```

Converting the preprocessed dataframe to CSV file for Tableau Visualization.


Non-Diabetic vs Diabetic



Age Bins and Avg BP



Age (bin) 2



Distribution Of Diabetic And Non Diabetic

